

How to Teach a 9B Model Your Chip

ARA Labs

Frontier language models do not know your internal code. On a 125-question closed-book benchmark about a real hardware platform, spanning register maps, memory maps, build pins, driver conventions, and issue history, Claude Opus 5 answers 72.0%; a 9B open-weight model that we adapt on the platform’s own repositories answers **92.8%**. Getting there is not a matter of running continued pretraining on the internal corpus: that obvious recipe *silently fails*, halving training loss while leaving recall of memorized facts exactly unchanged. This report documents the decision chain of a fully open study on the PULP Carfield/Cheshire platform: which testbed to use (and why a popular open-source chip cannot serve as one), continued pretraining versus fine-tuning versus retrieval, and, centrally, a knowledge-rewriting augmentation stage whose design rules (full coverage of all facts, template diversity over repetition, reversed statements, whole-table narratives against similar-fact interference) each trace to a measured failure of a simpler variant. We also report how our own auto-generated benchmark misled us twice, through lexically leaky multiple-choice questions and format-incomplete few-shot exemplars, and how we repaired it. All code, the data pipeline, the benchmark, and the model weights are released.

Code: github.com/ARA-Labs/PULP-LLM

Model: Qwen3.5-9B-PULP-DAPT (Hugging Face)



1. Introduction

Closed-source frontier models do not know how your internal systems work. The gap is not syntax: today’s models write fluent Verilog and C. What they lack is everything your organization never published, and that is most of what an internal assistant actually needs: which register sits at which offset, which memory region an address belongs to, how your packages depend on each other, which driver function configures which block, why a design decision was made, and what a two-line note in a register description implies. This tacit, organization-internal knowledge is increasingly recognized as the binding constraint on LLM assistance for real software and hardware work (Peng et al., 2025), and no amount of prompting recovers facts that were never in the training set. An assistant that guesses a register offset is worse than no assistant at all.

So: can a team *train* a model that knows its system? Chip-NeMo (Liu et al., 2023) showed at NVIDIA that domain-adaptive pretraining (DAPT) over internal data can work, but its corpus, evaluations, and most operational details are internal, its compute budget is beyond a small team, and the public account leaves the practitioner’s questions open: what do you do first, what fails, and how do you know your model *knows* rather than retrieves or guesses?

This report answers those questions with a small, fully reproducible study. We picked a real, actively developed but obscure hardware platform, PULP Carfield/Cheshire (ETH Zürich), as a public stand-in for a proprietary chip, and asked whether knowledge injection into an open 9B model

can be made to work and be shown to work (Figure 1). The answer is yes: on a 125-question layered closed-book benchmark, the adapted model answers 92.8% versus 72.0% for Claude Opus 5 and 43.2% for its own untrained base (Figure 2); a held-out audit over the entire auto-generated question bank confirms the gap is not an artifact of question selection.

The number, however, is not the contribution. The contribution is the *decision chain*: at every fork of this project there were several plausible options, the default option usually failed, and the failure was measurable. Continued pretraining on the raw corpus, the recipe most teams would try first, reduced training loss from 0.80 to 0.35 while leaving memorization accuracy *exactly unchanged*, an industrial-scale confirmation of the extraction results of Allen-Zhu and Li (2023). Augmenting only the facts we happened to quiz inflated the score on those facts and did nothing anywhere else. Our own auto-generated benchmark twice reported numbers that had nothing to do with knowledge, once through lexical leakage in multiple-choice construction and once through a missing answer-format exemplar in the few-shot prompt. Each of these failures produced a design rule, and the final recipe is nothing more than the accumulation of those rules.

Concretely, this report documents:

- **Problem framing** (§2): why “knowledge in weights” is the requirement worth testing, and why public code-generation benchmarks, which frontier models now saturate, cannot differentiate an internal-knowledge model.




Figure 1. Overview. A chip platform’s knowledge (register maps, memory maps, drivers, build system, issue history) exists in no pretraining corpus; we use the open PULP Carfield SoC as a stand-in for a proprietary chip. Its 259 repositories become a 31.2M-token corpus plus a 3.5M-token *knowledge-rewriting augmentation* set (every extracted fact restated through 24 diverse templates, one third reversed, plus whole-table narrative documents), since continued pretraining on the raw corpus alone injects no retrievable facts (§5). After 77 minutes of full-parameter continued pretraining on 4×H100, the model is examined closed-book on 1,776 auto-generated, machine-checked questions, where it answers 81.6% versus 72.2% for Claude Opus 5 and 41.2% for its untrained base (full-bank numbers; §4).

- **Design decisions** (§3): six forks in the road (testbed, adaptation method, base model, data recipe, evaluation design, systems configuration), each presented as the options we considered, the choice we made, and the measurement behind it.
- **Results** (§4): the ablation ladder from 43.2% to 92.8%, layer-by-layer profiles against a frontier baseline, and a held-out anti-overfitting check on never-inspected questions.
- **Findings** (§5): eight transferable lessons on knowledge injection and on auditing auto-generated benchmarks, with the evidence for each.

Everything is open: the corpus pipeline, the augmentation generator, the benchmark and its generator, per-question results for every model, the trained weights, and the training configuration.¹

¹Code and benchmark: <https://github.com/ARA-Labs/PULP-LLM>. Weights: <https://huggingface.co/AgentNativeResearchLab/Qwen3.5-9B-PULP-DAPT>.

2. The Problem We Are Actually Solving

The customer’s question. A chip team with a fully closed-source design wants an internal assistant that knows the chip the way a senior engineer does. Before committing to a training program, two prior questions must be answered with evidence rather than enthusiasm.

Question 1: where can a self-trained model beat the frontier? Not on public skills. We reproduced the VeriThoughts fine-tuning pipeline (Yubeaton et al., 2025) as a calibration exercise and then ran Claude Opus 5 on the same formally-verified Verilog benchmark: the frontier model scores 95.3% pass@1, above the fine-tuned specialist models the literature celebrates. Public RTL generation is saturating; a 7–9B model fine-tuned on public data differentiates nothing. What a frontier model *cannot* have is knowledge that never reached its training set: the proprietary chip itself. Knowledge injection is therefore not one use case among many; it is the only battleground where the self-trained model holds structural advantage.




Figure 2. The recipe, not the corpus, carries the result. Closed-book accuracy on the 125-question layered audit set. Continued pretraining on the raw 31M-token corpus leaves memorization at base level despite halving training loss (pilot measurement, §3.4); augmenting only the quizzed facts helps only those facts (55.2%); the full recipe (coverage of all facts, 24 diverse templates, reversed forms, whole-table narratives) reaches 92.8%, surpassing Claude Opus 5 (dashed line).

Question 2: what counts as “knowing”? We require *closed-book* recall: no retrieval, no context documents, the fact must come out of the weights. This is a deliberate, restrictive choice. Retrieval-augmented generation answers lookup questions well, but the engineer-facing failures we target (cross-file reasoning, “which function configures this register,” recall inside a debugging conversation where nobody pastes the register map) degrade precisely when knowledge lives only in an external index. Testing closed-book recall isolates the property we are trying to buy with training compute. (Whether RAG plus a frontier model is the better *product* for a given team is a separate question our open questions section returns to; this study measures what training can inject.)

Requirements. The study therefore needs: (i) a testbed whose facts are verifiably absent or scarce in frontier pre-training, yet real enough that results transfer to a customer engagement; (ii) an evaluation that is machine-checkable, layered across the kinds of knowledge an engineer uses, and hardened against the model answering by surface pattern rather than knowledge; (iii) a recipe whose every component is justified by an ablation, because the customer will ask “do we need this step?” for each one; and (iv) a footprint a small team can afford without procurement: every training run in this study fits on a single 4-GPU node in about an hour.

3. Design Decisions

Each subsection follows the same shape: the options on the table, the choice, and the measurement that forced it.

3.1. D1: The testbed. Why not a popular open-source chip

Options. (a) Use OpenTitan, the best-documented open SoC, as-is; (b) apply a consistent renaming/constant-perturbation transform to OpenTitan to synthesize a chip that provably exists nowhere (“Cisco-X”); (c) use a real but obscure platform.

Choice. (c), the PULP Carfield/Cheshire platform, with (b) validated and kept as a companion methodology.

Why. A knowledge-injection demo is meaningless if the model already knows the answers. We probed Claude Opus 5 closed-book on 252 auto-generated OpenTitan questions: 75.0% correct, including exact register offsets. The chip is *in the weights*, so option (a) cannot attribute anything to training. The Cisco-X transform (identifier renaming through the register-tool regeneration flow, constant perturbation, layout shuffling) drives frontier closed-book performance to 0.0% with 99.6% explicit abstention while open-book performance stays at 91.8%: clean attribution, and a reusable trick for customer engagements where even question texts are sensitive. But a synthetic chip cannot demonstrate handling of *real* corpus mess: Carfield is actively developed, sparsely discussed on the public web (Opus 5: 72.0% on our final benchmark, with register offsets at 18%), and comes with the full organic sprawl (259 repositories, generated headers, issues, drivers) that a customer’s internal data will have.

3.2. D2: Adaptation method. DAPT before SFT, weights before retrieval

Options. (a) Retrieval over the corpus with a strong general model; (b) supervised fine-tuning on Q/A pairs; (c) domain-adaptive continued pretraining (DAPT), optionally followed by SFT.

Choice. (c), DAPT first; SFT deferred until knowledge is in the weights.

Why. The property under test is closed-book knowledge (§2); retrieval by construction does not put knowledge in weights, so (a) answers a different question. SFT-first (b) teaches answer *format* but must invent its Q/A pairs from the same corpus, and our corpus is 31.2M tokens of mostly non-Q/A material; at that scale the ChipNeMo-style recipe (Liu et al., 2023) (full-parameter continued pretraining at low learning rate with a general-data replay mix) is the established path, and our token budget sits comfortably in the regime where continued pretraining is worthwhile. The 3-shot completion protocol substitutes for instruction-following during evaluation, letting us measure injected knowledge without an SFT stage confounding it.

3.3. D3: Base model. A base checkpoint, not a chat model

Options. (a) Instruction-tuned 7–9B checkpoints; (b) base checkpoints.

Choice. Qwen3.5-9B-Base, full-parameter, LR 5×10^{-6} cosine, packing at 4096, two epochs.

Why. ChipNeMo’s ablations are unambiguous: DAPT on a chat model degrades both knowledge uptake and chat ability, and low learning rate with full-parameter updates (not LoRA) is what makes fact absorption work. We saw no reason to relitigate them on our budget. We took the newest open base model at the largest size our 4-GPU budget trains in about an hour.

3.4. D4: The data recipe. Where the result actually comes from

This is the decision that produced the paper’s central lesson, in three measured steps.

Step 1: raw corpus fails. The obvious recipe, feeding the filtered 31.2M-token organizational corpus (RTL, generated headers, docs, drivers, issues of all 259 non-fork repositories, plus 8% general-text replay) to the trainer, *looks* like it works: training loss falls from 0.80 to 0.35. It injects nothing extractable: on the pilot benchmark’s 42 register-offset questions the raw-DAPT model scores 1/42, identical to base. A fact that appears in exactly one surface form (a `#define` line) is stored but not retrievable as an answer, exactly as [Allen-Zhu and Li \(2023\)](#) predict with synthetic biographies; we believe this is the clearest industrial-corpus confirmation of that result to date, and it means the popular mental model “CPT on your data $\Rightarrow$ your model knows your data” fails *silently*: loss curves report success while extraction stays at zero.

Step 2: augmentation works exactly as far as it covers.

We then restated facts in multiple declarative paraphrases mixed into the corpus. The first version covered only the 461 facts our pilot questions quizzed, and its gains land *only* on layers those facts populate. On the final benchmark, that model (55.2%) matches the untrained base *exactly* on everything its augmentation never covered: dependency pins 0/8, region sizes 1/8, driver–register links 4/15. Coverage determines the extractable scope; evaluation must therefore be a random audit of the full fact inventory, never the augmentation’s target list.

Step 3: diversity, reversal, and context against interference. The final generator enumerates *every* fact from every structured source (~ 340 register offsets, all register-description fields, both memory maps, build pins, driver–register links from static analysis, issue titles) and emits each through 24 LLM-written templates per fact type (gen-

erated with ten LLM calls, validated programmatically), of which at least a third are *reversed* (value first, entity second) to dodge the reversal curse ([Berglund et al., 2023](#)). Two further rules come from an interference observation: at identical exposure, 9 mutually distinctive base addresses were memorized 8/9 while 339 near-identical offsets managed 9/42. Similarity, not exposure count, was the bottleneck. The countermeasures are whole-table narrative documents (the register map rendered as ascending prose, a markdown table, a C header, and a neighbor-chain walk, so that similar facts appear as each other’s context) and fewer repeats (6 rather than 8) in favor of more forms. The augmentation totals 3.49M tokens, about 10% of the training mix. Result: register-offset recall went from 1/42 (raw) and 9/42 (quizzed-facts augmentation) to 25/42 on the pilot protocol, and to 97% (323/333) over *all* offsets under the final protocol.

3.5. D5: Evaluation. Auditing our own benchmark

Options. (a) Quiz whatever is easy to auto-generate; (b) design the distribution top-down from what an engineer needs to know, and treat the generator itself as a system under test.

Choice. (b), after (a) burned us.

Why. Our pilot benchmark drew 54% of its questions from register offsets, not because offsets matter most, but because a regex over generated headers made them free. Auto-generated evaluations inherit the streetlight effect of their sources, and training data designed against such an evaluation inherits it twice. The final benchmark is layered by knowledge kind: L1 facts, 30%; L2 structure (memory-map ownership, sizes, dependency pins), 20%; L3 behavior (documentation semantics as 4-way multiple choice), 25%; L4 cross-source (which driver function touches which register, from static analysis), 15%; L5 engineering history, 10%. Every question is machine-checkable.

Three generator defects surfaced only because we ran a frontier model as an adversarial probe, and each became a construction rule. *Lexical leak*: Opus scored 100% on description $\rightarrow$ register-name multiple choice by word overlap alone (“Ethernet RGMII PHY clock divider” $\rightarrow$ ETH_RGMII_PHY_CLK_DIV); we now keep a multiple-choice question only if every distractor’s lexical-overlap score is at least the answer’s, and we deleted one question type (RTL instance $\rightarrow$ module) that could not be repaired. *Few-shot format holes*: a shot set whose three exemplars were all bit-range answers caused models to answer hex questions in decimal, mis-scored by the normalizer *and* genuinely degraded, silently depressing an entire protocol; every answer format now has an exemplar. *Worthless question types*: issue-number $\rightarrow$ title recall is a pure arbitrary mapping with no engineering value; no model exceeded 27%, and we removed the type (1,423 questions) from the

Table 1. Closed-book accuracy on the 125-question layered audit set (3-shot completion, greedy decoding; letters for multiple choice, exact match after normalization otherwise). “Quizzed-facts aug.” is the ablation whose augmentation covered only the 461 pilot-quizzed facts.

Model	Total	L1 ₄₀	L2 ₂₆	L3 ₃₂	L4 ₂₀	L5 ₇
Qwen3.5-9B-Base	43.2%	7	4	29	11	3
+ quizzed-facts aug.	55.2%	20	4	31	9	5
Claude Opus 5	72.0%	23	18	32	11	6
DAPT (ours)	92.8%	40	19	32	20	5

Table 2. Full bank: all 1,776 auto-generated questions, of which 1,651 were never inspected during development.

Model	Total	L1 ₄₁₉	L2 ₅₃	L3 ₅₃	L4 ₂₅	L5 ₁₂₂₆
Qwen3.5-9B-Base	41.2%	7%	32%	89%	52%	51%
Claude Opus 5	72.2%	28%	79%	98%	52%	86%
DAPT (ours)	81.6%	97%	68%	100%	100%	76%

bank rather than let it dominate the aggregate.

3.6. D6: Systems configuration. The cheap 2.2×

Three configuration-level choices, found by 15-step smoke tests, cut cost without touching the science: the `flash-linear-attention` kernels are mandatory for Qwen3.5’s gated-delta-net layers (25 s/it → 7.4 s/it); trading ZeRO-3’s per-layer parameter gathering for either ZeRO-2 or a micro-batch of 2 (Rajbhandari et al., 2019) gives an equivalent further 2.2× (7.43 → 3.41 s/it), making the final training run 77 minutes on 4×H100; and evaluation under vLLM (Kwon et al., 2023) with prefix caching amortizes the shared few-shot prompt across the whole bank. Training used LLaMA-Factory (Zheng et al., 2024) with DeepSpeed; the allocator setting `expandable_segments` suppressed fragmentation warnings that otherwise threaten long runs.

4. Results

Audit set. Table 1 and Figure 2: the full recipe reaches 92.8%, +20.8 points over Claude Opus 5 and +49.6 over its own base. The quizzed-facts ablation row is the coverage result of §3.4 in tabular form: +13 questions on covered layers, base-level everywhere else.

Anti-overfitting check. Because the audit set is the only part of the benchmark we ever looked at while iterating, the full bank (Table 2) is the honest test: on the 379 never-inspected L1 fact questions the model scores 97%, matching its 100% on the inspected 40. The subset number is not sample luck, and the augmentation’s full-coverage property is real.

Complementary failure profiles. The layer columns of Table 2 show why “beats Opus” is the right summary but not the interesting one. The frontier model wins wherever semantic association helps: issue-title → repository attribution (86%), memory-region ownership (26/27), documentation semantics (98%). It collapses on pure memorization: 18% on register offsets. The adapted 9B is the mirror image: 97% on offsets, 100% on driver-register links, weaker on numeric range reasoning (region ownership 11/27). Knowledge injection buys exactly the layer that retrieval-free internal assistance needs and that scale alone does not provide.

Compute. Every training run in this study is a single 4×H100 node for 77 minutes; evaluation of the full bank is one GPU-hour under vLLM.

5. Findings

Eight lessons, each traceable to a measurement above; we believe they transfer to any “teach a model our internal system” effort.

- 1. Raw-corpus continued pretraining fails silently.** Loss halves; extraction does not move (1/42 → 1/42). Never accept a loss curve as evidence of knowledge injection.
- 2. Knowledge rewriting is the enabling stage, not an optimization.** Every point of closed-book gain in this study is downstream of restating facts in many surface forms. Corollary for data cleaning: redundant restatements of the same fact in a customer corpus are *assets*; semantic deduplication destroys exactly the diversity that makes facts extractable.
- 3. Coverage determines the extractable scope.** Augmented layers improve; unaugmented layers match the untrained base to the question. Enumerate the full fact inventory; let the benchmark sample it.
- 4. Diversity beats repetition, and word order is part of diversity.** 24 templates × 6 repeats beats 6 templates × 8; reversed statements are required, not optional (Berglund et al., 2023).
- 5. Interference, not exposure, limits dense fact tables.** At equal exposure: distinctive facts 8/9, near-identical facts 9/42. Give similar facts distinguishing context (whole-table narratives) rather than more repeats.
- 6. Auto-generated benchmarks need adversarial audits.** A frontier model probing our benchmark exposed lexical leaks (100% without knowledge) and format holes that a weaker model would have hidden. Budget for auditing the benchmark like code.

7. **Delete question types that measure nothing.** Arbitrary-mapping recall (issue number $\leftrightarrow$ title) stayed $\leq 27\%$ for every model and would have dominated the aggregate; a benchmark’s distribution is a design decision, not a harvest.
8. **The systems margin is real and cheap.** Kernel choice and parallelism configuration bought $7.6\times$ end-to-end training throughput in two 15-step smoke tests; at small budgets, configuration *is* capacity.

6. Related Work

ChipNeMo (Liu et al., 2023) established the DAPT-for-chip-knowledge paradigm at industrial scale and supplies several rules we adopt unmodified (base-model start, full-parameter low-LR training, replay mixing); our study differs in being fully open, three orders of magnitude cheaper, and centered on the question ChipNeMo’s internal evaluation could not expose publicly: *what makes injected knowledge extractable*. The answer traces to Allen-Zhu and Li (2023), who showed with synthetic corpora that facts stated in a single form resist question-answering extraction unless training data is augmented with paraphrases; we reproduce this on a real 31M-token engineering corpus and quantify its coverage boundary. The reversal curse (Berglund et al., 2023) motivates our reversed templates. On the evaluation side, code-generation benchmarks with formal verification (Yubeaton et al., 2025) measure a skill frontier models now largely possess (95.3% pass@1 in our reproduction); our benchmark instead targets closed-book platform knowledge, where the same frontier model scores 72%. Training and serving infrastructure: LLaMA-Factory (Zheng et al., 2024), ZeRO (Rajbhandari et al., 2019), vLLM (Kwon et al., 2023).

7. Limitations and Open Questions

Stated plainly, because each one shapes what we would do next.

No forgetting audit yet. We have not measured general-capability regression (e.g. an MMLU subset) after DAPT; the 8% replay mix follows ChipNeMo practice but is unverified here. This is the first gap we would close.

Single seed, single platform. Every number is one training run on one platform; the ablation ladder’s steps are large (+37.6 points for the recipe change) but error bars would still become necessary before any causal claim about smaller effects.

Bundled recipe components. Full coverage, template count, reversal, and whole-table narratives shipped together in the final run; their individual contributions are not isolated. The interference explanation (Finding 5) in par-

ticular rests on one observation and needs a controlled similarity $\times$ exposure $\times$ context experiment.

L3 does not differentiate. Register naming is descriptive by design, so description $\leftrightarrow$ name multiple choice is partially solvable by semantic matching; all models score $\geq 89\%$ there even after leak filtering.

Remaining model weaknesses. Numeric range membership (which region contains address X : 11/27) resists declarative augmentation; it needs reasoning over the map, suggesting either chain-style augmentation or an SFT stage; issue-history recall remains weak for all models.

No RAG comparison. This study measures what training injects into weights; a deployment-oriented comparison against retrieval (and against retrieval *plus* our adapted model) is the obvious next experiment, and we expect the hybrid to win.

References

- Zeyuan Allen-Zhu and Yuanzhi Li. Physics of language models: Part 3.1, knowledge storage and extraction. *arXiv preprint arXiv:2309.14316*, 2023.
- Lukas Berglund, Meg Tong, Max Kaufmann, Mikita Balesni, Asa Cooper Stickland, Tomek Korbak, and Owain Evans. The reversal curse: LLMs trained on “A is B” fail to learn “B is A”. *arXiv preprint arXiv:2309.12288*, 2023.
- Woosuk Kwon et al. Efficient memory management for large language model serving with pagedattention. *arXiv preprint arXiv:2309.06180*, 2023.
- Mingjie Liu et al. Chipnemo: Domain-adapted llms for chip design. *arXiv preprint arXiv:2311.00176*, 2023.
- Xin Peng et al. Code digital twin: Empowering llms with tacit knowledge for complex software development. *arXiv preprint arXiv:2510.16395*, 2025.
- Samyam Rajbhandari, Jeff Rasley, Olatunji Ruwase, and Yuxiong He. Zero: Memory optimizations toward training trillion parameter models. *arXiv preprint arXiv:1910.02054*, 2019.
- Patrick Yubeaton et al. Verithoughts: Enabling automated verilog code generation using reasoning and formal verification. *arXiv preprint arXiv:2505.20302*, 2025.
- Yaowei Zheng et al. Llamafactory: Unified efficient fine-tuning of 100+ language models. *arXiv preprint arXiv:2403.13372*, 2024.

A. Reproduction

The repository (<https://github.com/ARA-Labs/PULP-LLM>) contains the corpus crawler, fact extractors, template generator, augmentation builder, training and evaluation scripts, the full benchmark with per-model per-question results. The end-to-end sequence is: crawl the organization ($\rightarrow$ 31.2M tokens after filtering toolchain forks and test payloads); generate the question bank and augmentation; train (stage=pt, full-parameter, LR 5×10^{-6} cosine, global batch 16, packing at 4096, bf16, 2 epochs, ZeRO-3 with micro-batch 2, $4 \times H100$, 77 min); evaluate with vLLM (3-shot completion, greedy, 24 max tokens, prefix caching). Model weights: <https://huggingface.co/AgentNativeResearchLab/Qwen3.5-9B-PULP-DAPT>.

B. Engineering Notes

Notes that cost us time and are cheap to pass on:

(i) Qwen3.5's gated-delta-net layers train $3.4\times$ slower without the `flash-linear-attention` package (the fallback path emits no warning); (ii) ZeRO-2 and ZeRO-3 with micro-batch 2 are throughput-equivalent here (3.47 vs. 3.41 s/it), so pick by memory headroom; (iii) `PYTORCH_CUDA_ALLOC_CONF=expandable_segments:true` silences per-step allocator retry warnings on tight configurations; (iv) organizational corpora hide enormous toolchain forks (ours: 168M characters of `binutils/glibc` forks, $5\times$ the real content); filter by repository identity before tokenizing; (v) few-shot exemplars must cover every answer format in the bank (§3.5); (vi) keep per-question model outputs forever: every re-scoring in this study (question deletions, protocol audits) was free because raw outputs were retained.